

Design Document : Memory Management Simulator

Name: Dara Zenas Zephaniah

Enrollment: 23321010

① Memory Layout and Assumptions - The Simulator models the

computer's physical memory as a contiguous block of raw bytes, managed dynamically at runtime.

Physical memory size : 1024 Bytes (configurable).

Addressing : Byte-addressable simulation.

Data Structure: The memory is not tracked by a bitmap but by a Double

Linked List of Memory Segments structures.

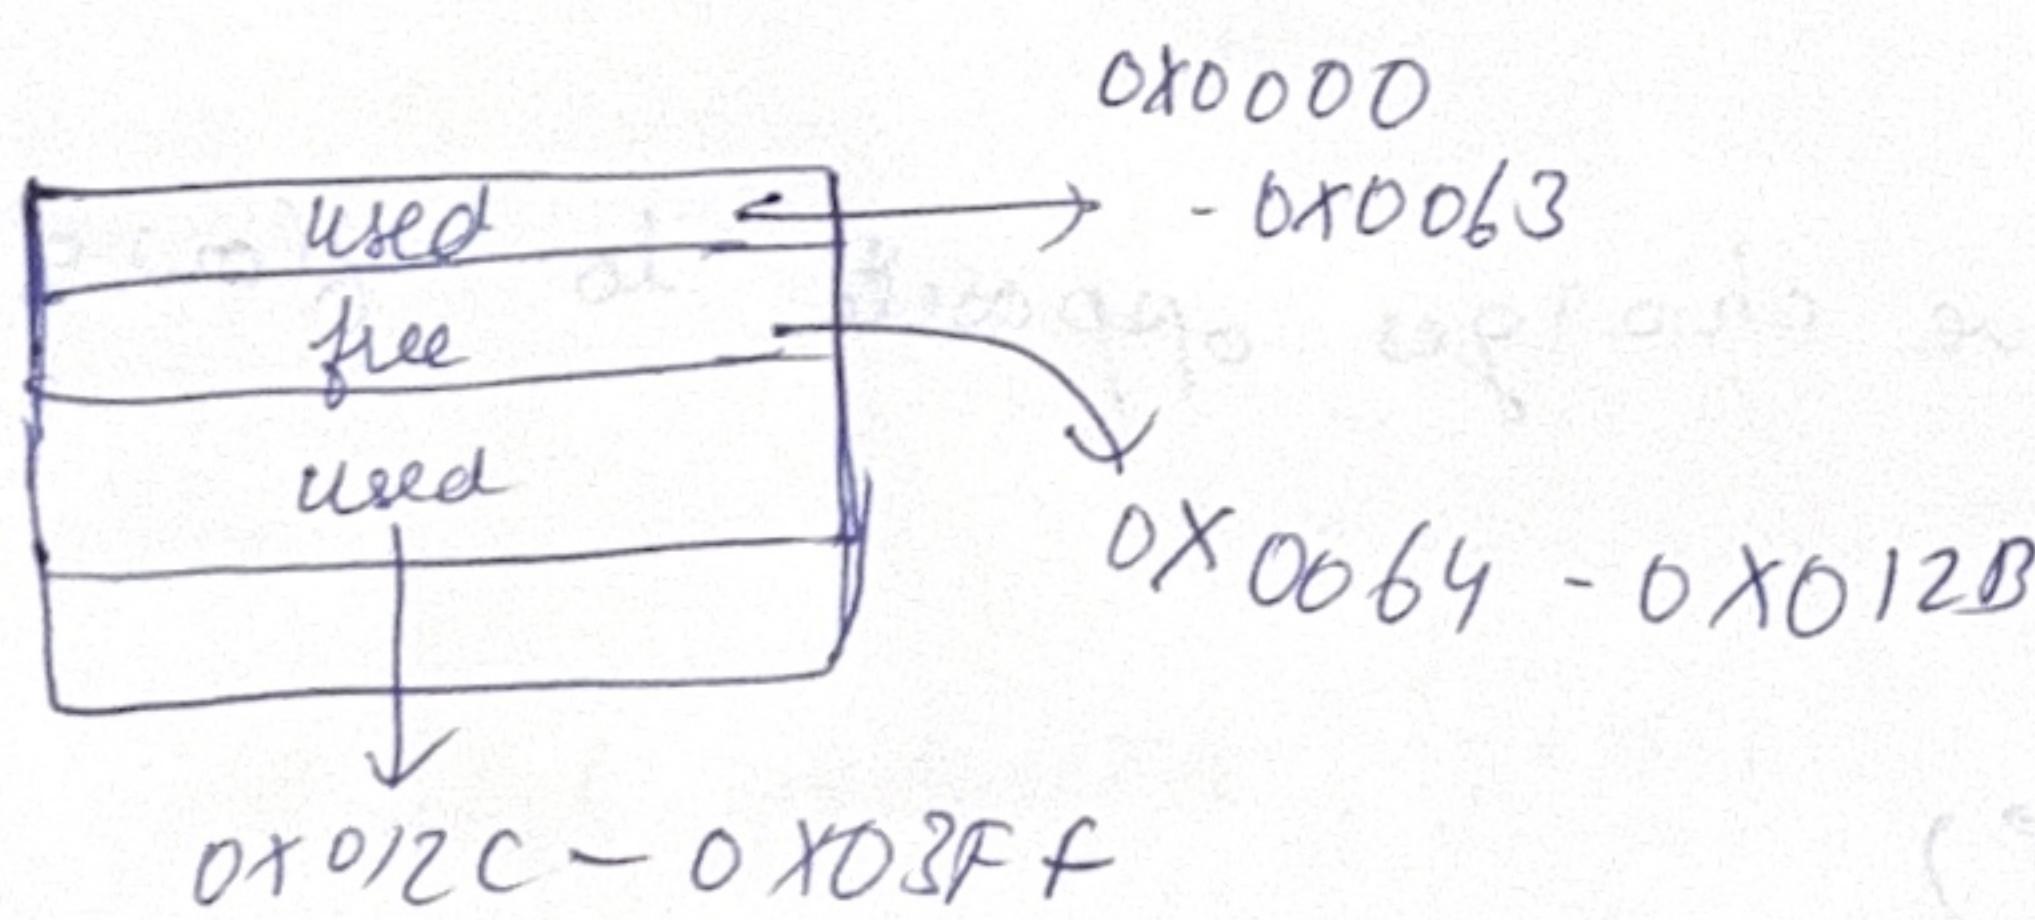
Metadata: Each segment header stores

start_addr: Physical starting address

size: Size of the block in bytes

is-allocated: Boolean flag (True = Used, False = Free).

id: Unique identifier for tracking allocation (0 to 1023).



② Allocation Strategy Implementation

Algorithm: First Fit

The Simulator implements the first fit allocation strategy, prioritized for its speed and simplicity.

a) Search: The allocator traverses the linked list from the head (start_addr: 0)

b) Selection: It selects the first free block that satisfies

block_size ≥ requested_size

c) Splitting: If the selected block is larger than required, it is split into two:

Block A (Allocated): Shrinks to the requested size.

Block B (Free): Created from the remaining space and inserted into the list immediately after block A.

- d) Coalescing (Deallocation): When free(id) is called.
- The block is marked as ~~free~~ free.
 - The system checks the next and prev pointers.
 - If a neighbor is also free, they are merged into a single larger block to mitigate External ~~to~~ Fragmentation.

③ Cache Hierarchy and Replacement Policy - The system implements a Two-level (L1/L2) Inclusive Cache Hierarchy to reduce access latency.

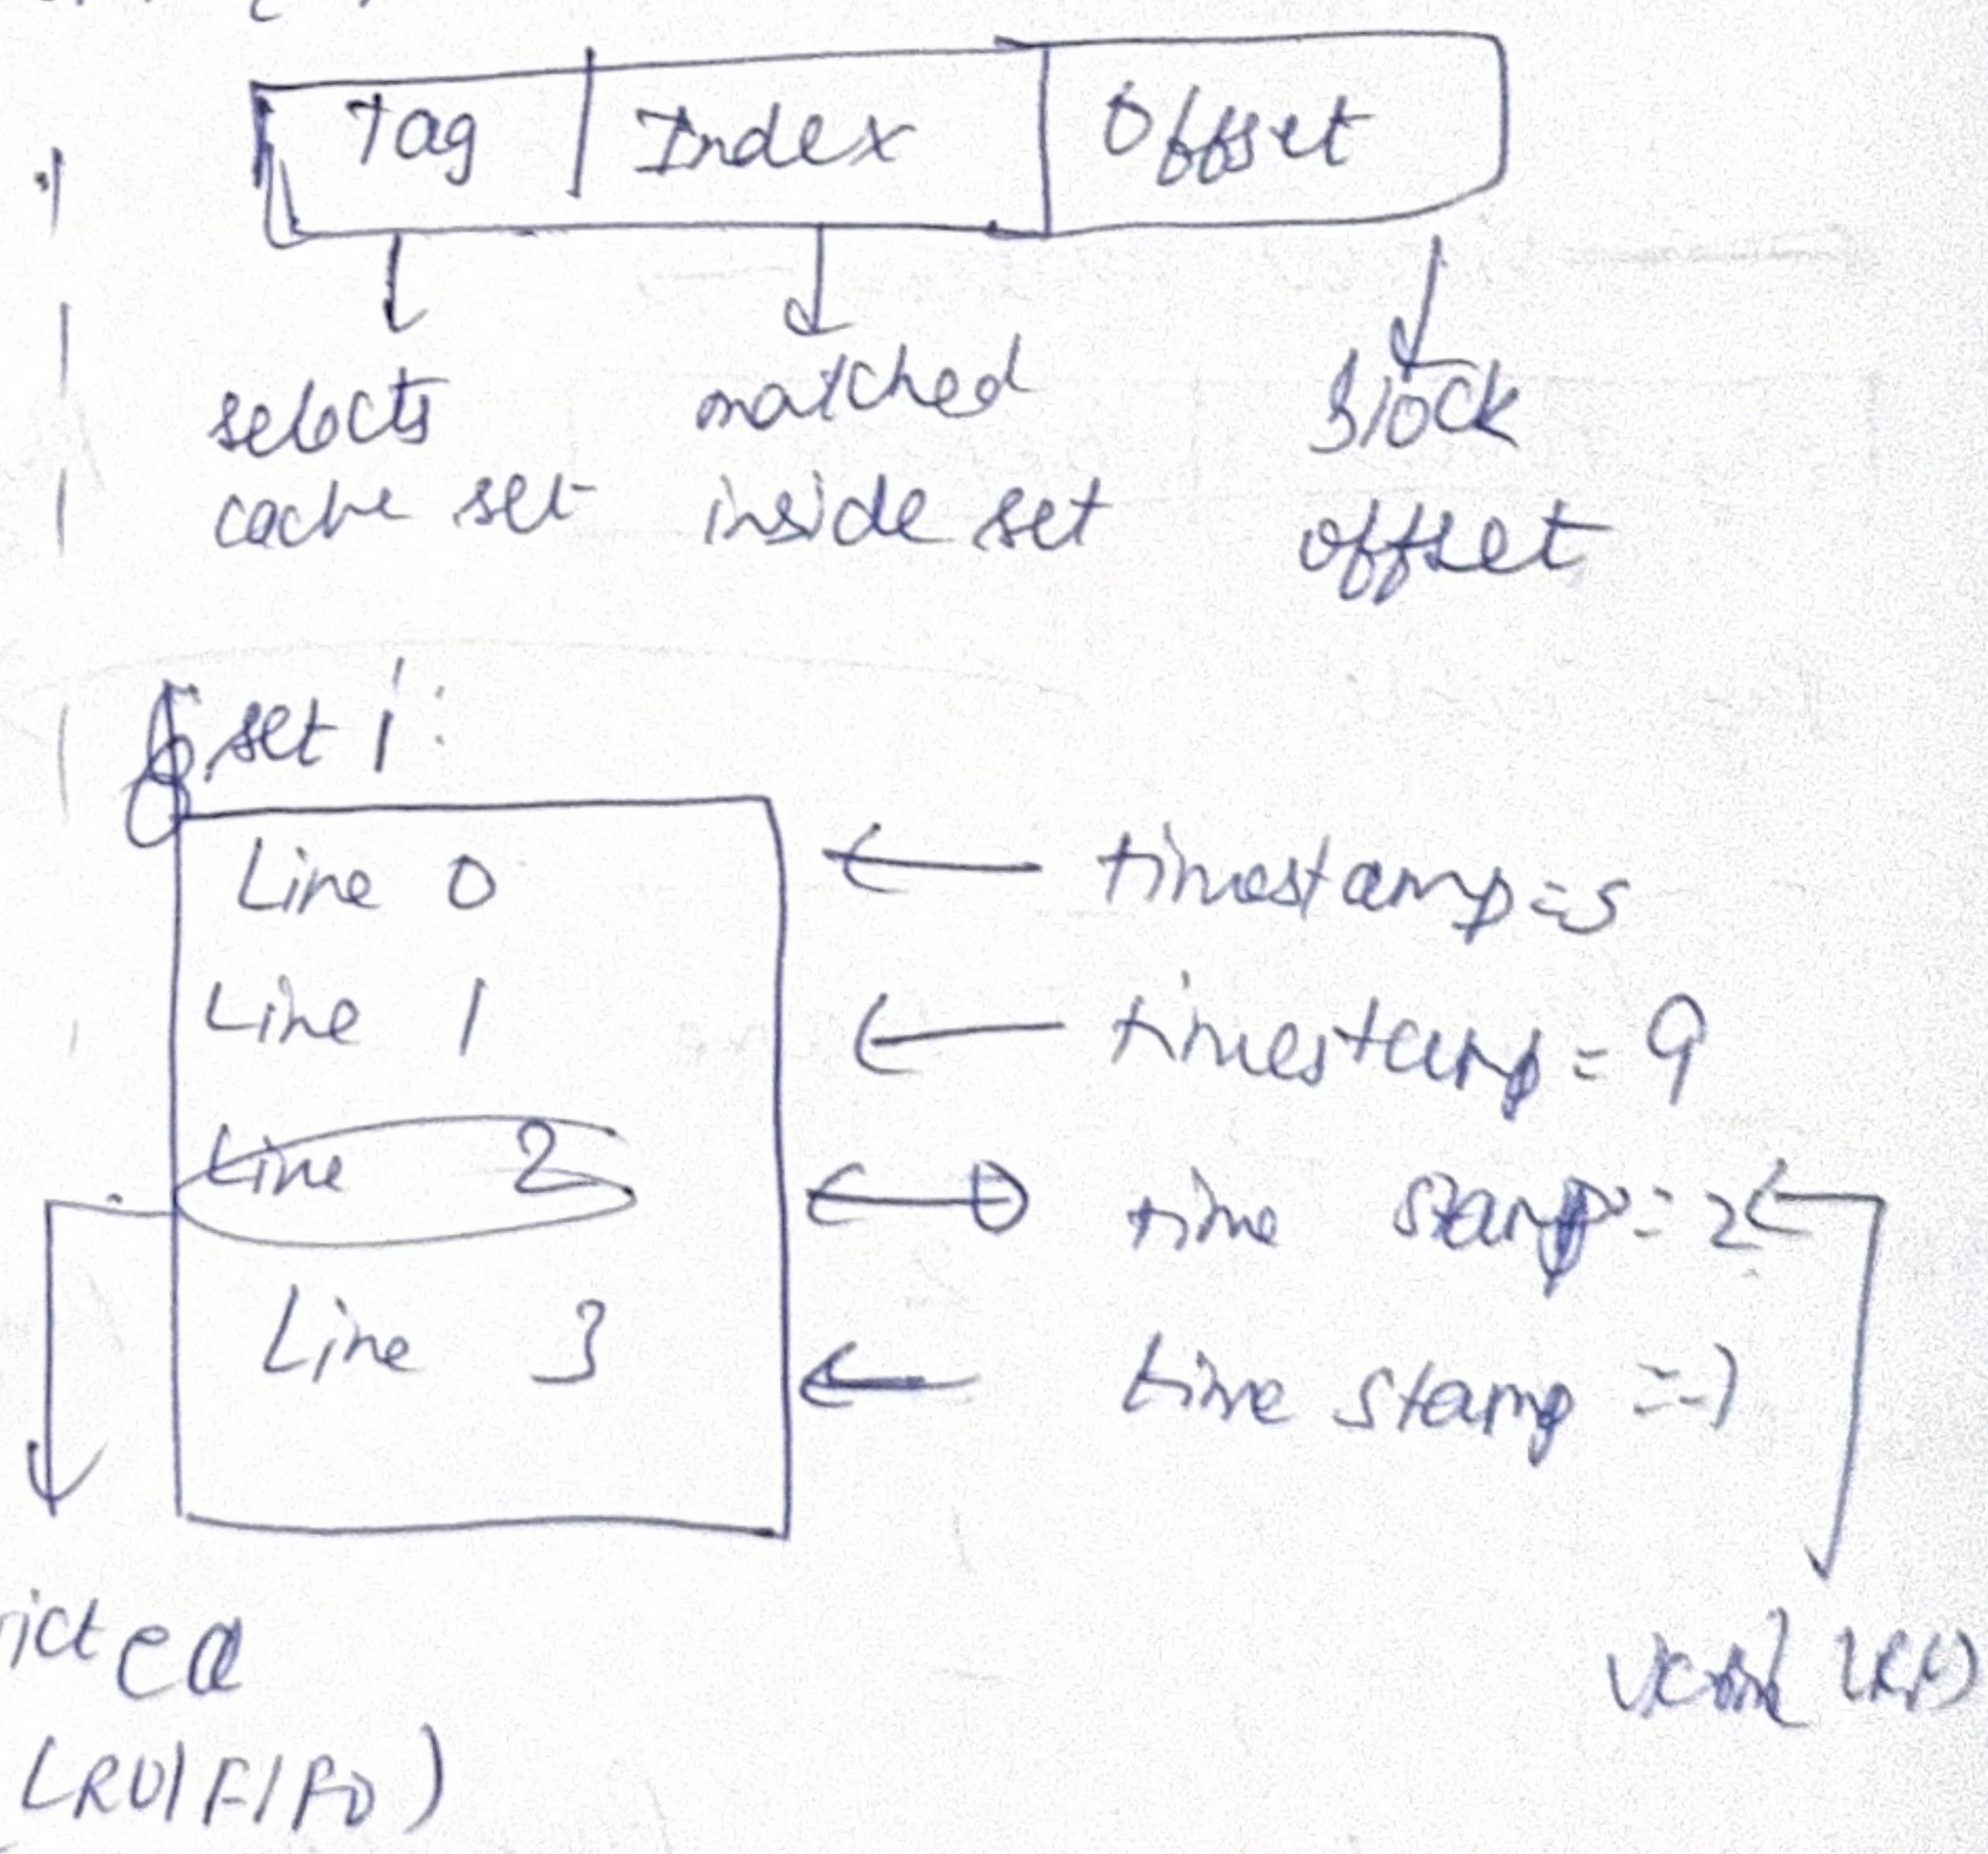
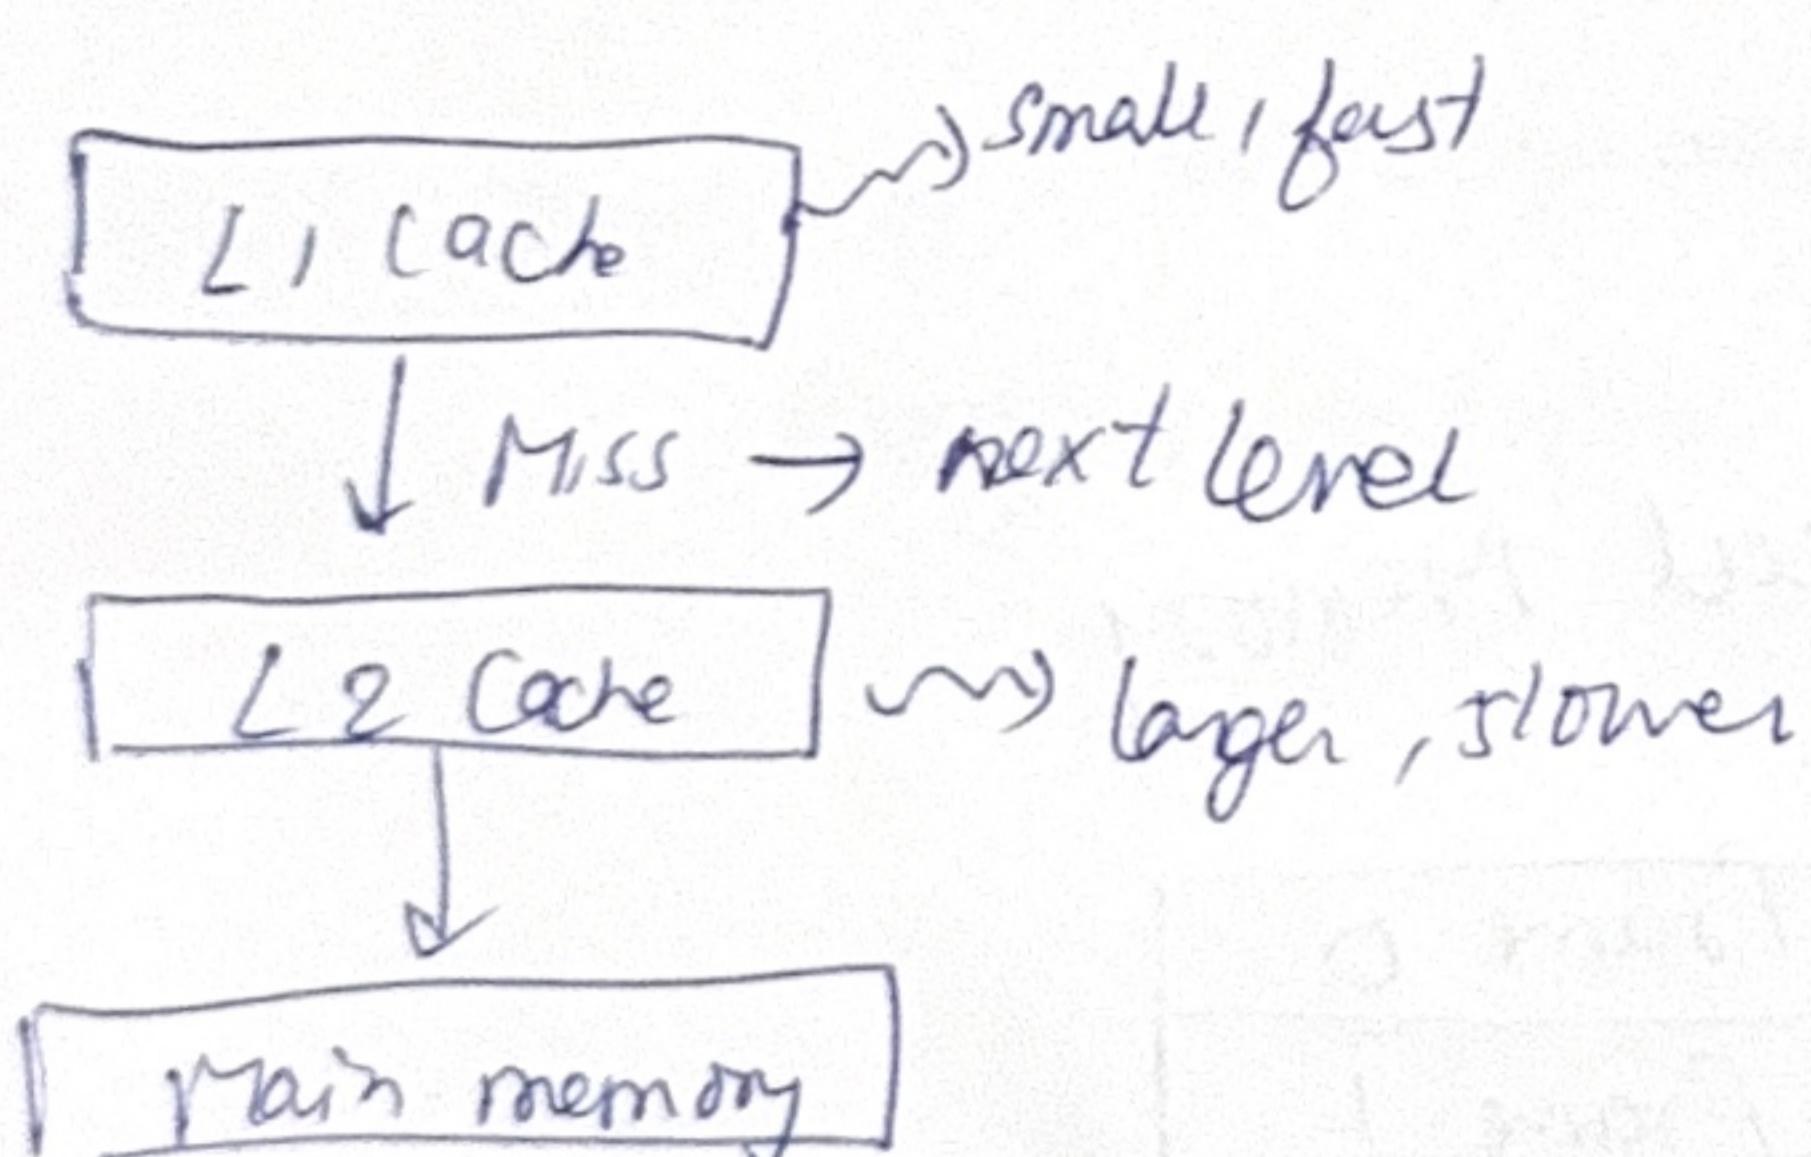
Cache specification

- L1 cache: 64 Bytes (4 lines \times 16 Byte Blocks). High speed, small capacity.
- L2 cache: 256 Bytes (16 lines \times 16 Byte Blocks). Slow speed, larger capacity.
- Mapping Policy: Direct mapped.
 - Each memory address maps to exactly one specific line in the cache based on the index bits.
 - formula Cache Line Index = $(\text{Address} / \text{block size}) \% \text{TotalLines}$

Address Decomposition (Bitwise Logic) - For a 16-Byte Block size:

- Offset: Lower 4 bits ($\log_2 16$).
- Index: Middle bits determined by line count (2 bits for L1, 4 bits for L2).
- Tag: Remaining upper bits used to verify identity.

Address format: TAG | INDEX | OFFSET



Access Flow

- ①) check L1: Extract L1 Index/Tag. If valid bit is set and tags match \rightarrow L1 HIT (1 cycle)
- ②) check L2: If L1 Miss, extract L2 index/tag. If L2 HIT \rightarrow L2 HIT (+10 cycles). Promote block to L1.
- ③) fetch RAM: If L2 Miss \rightarrow Access Main Memory (+1000 cycles). Load block into L2, then into L1.

④ Virtual Memory Model (Simplification)

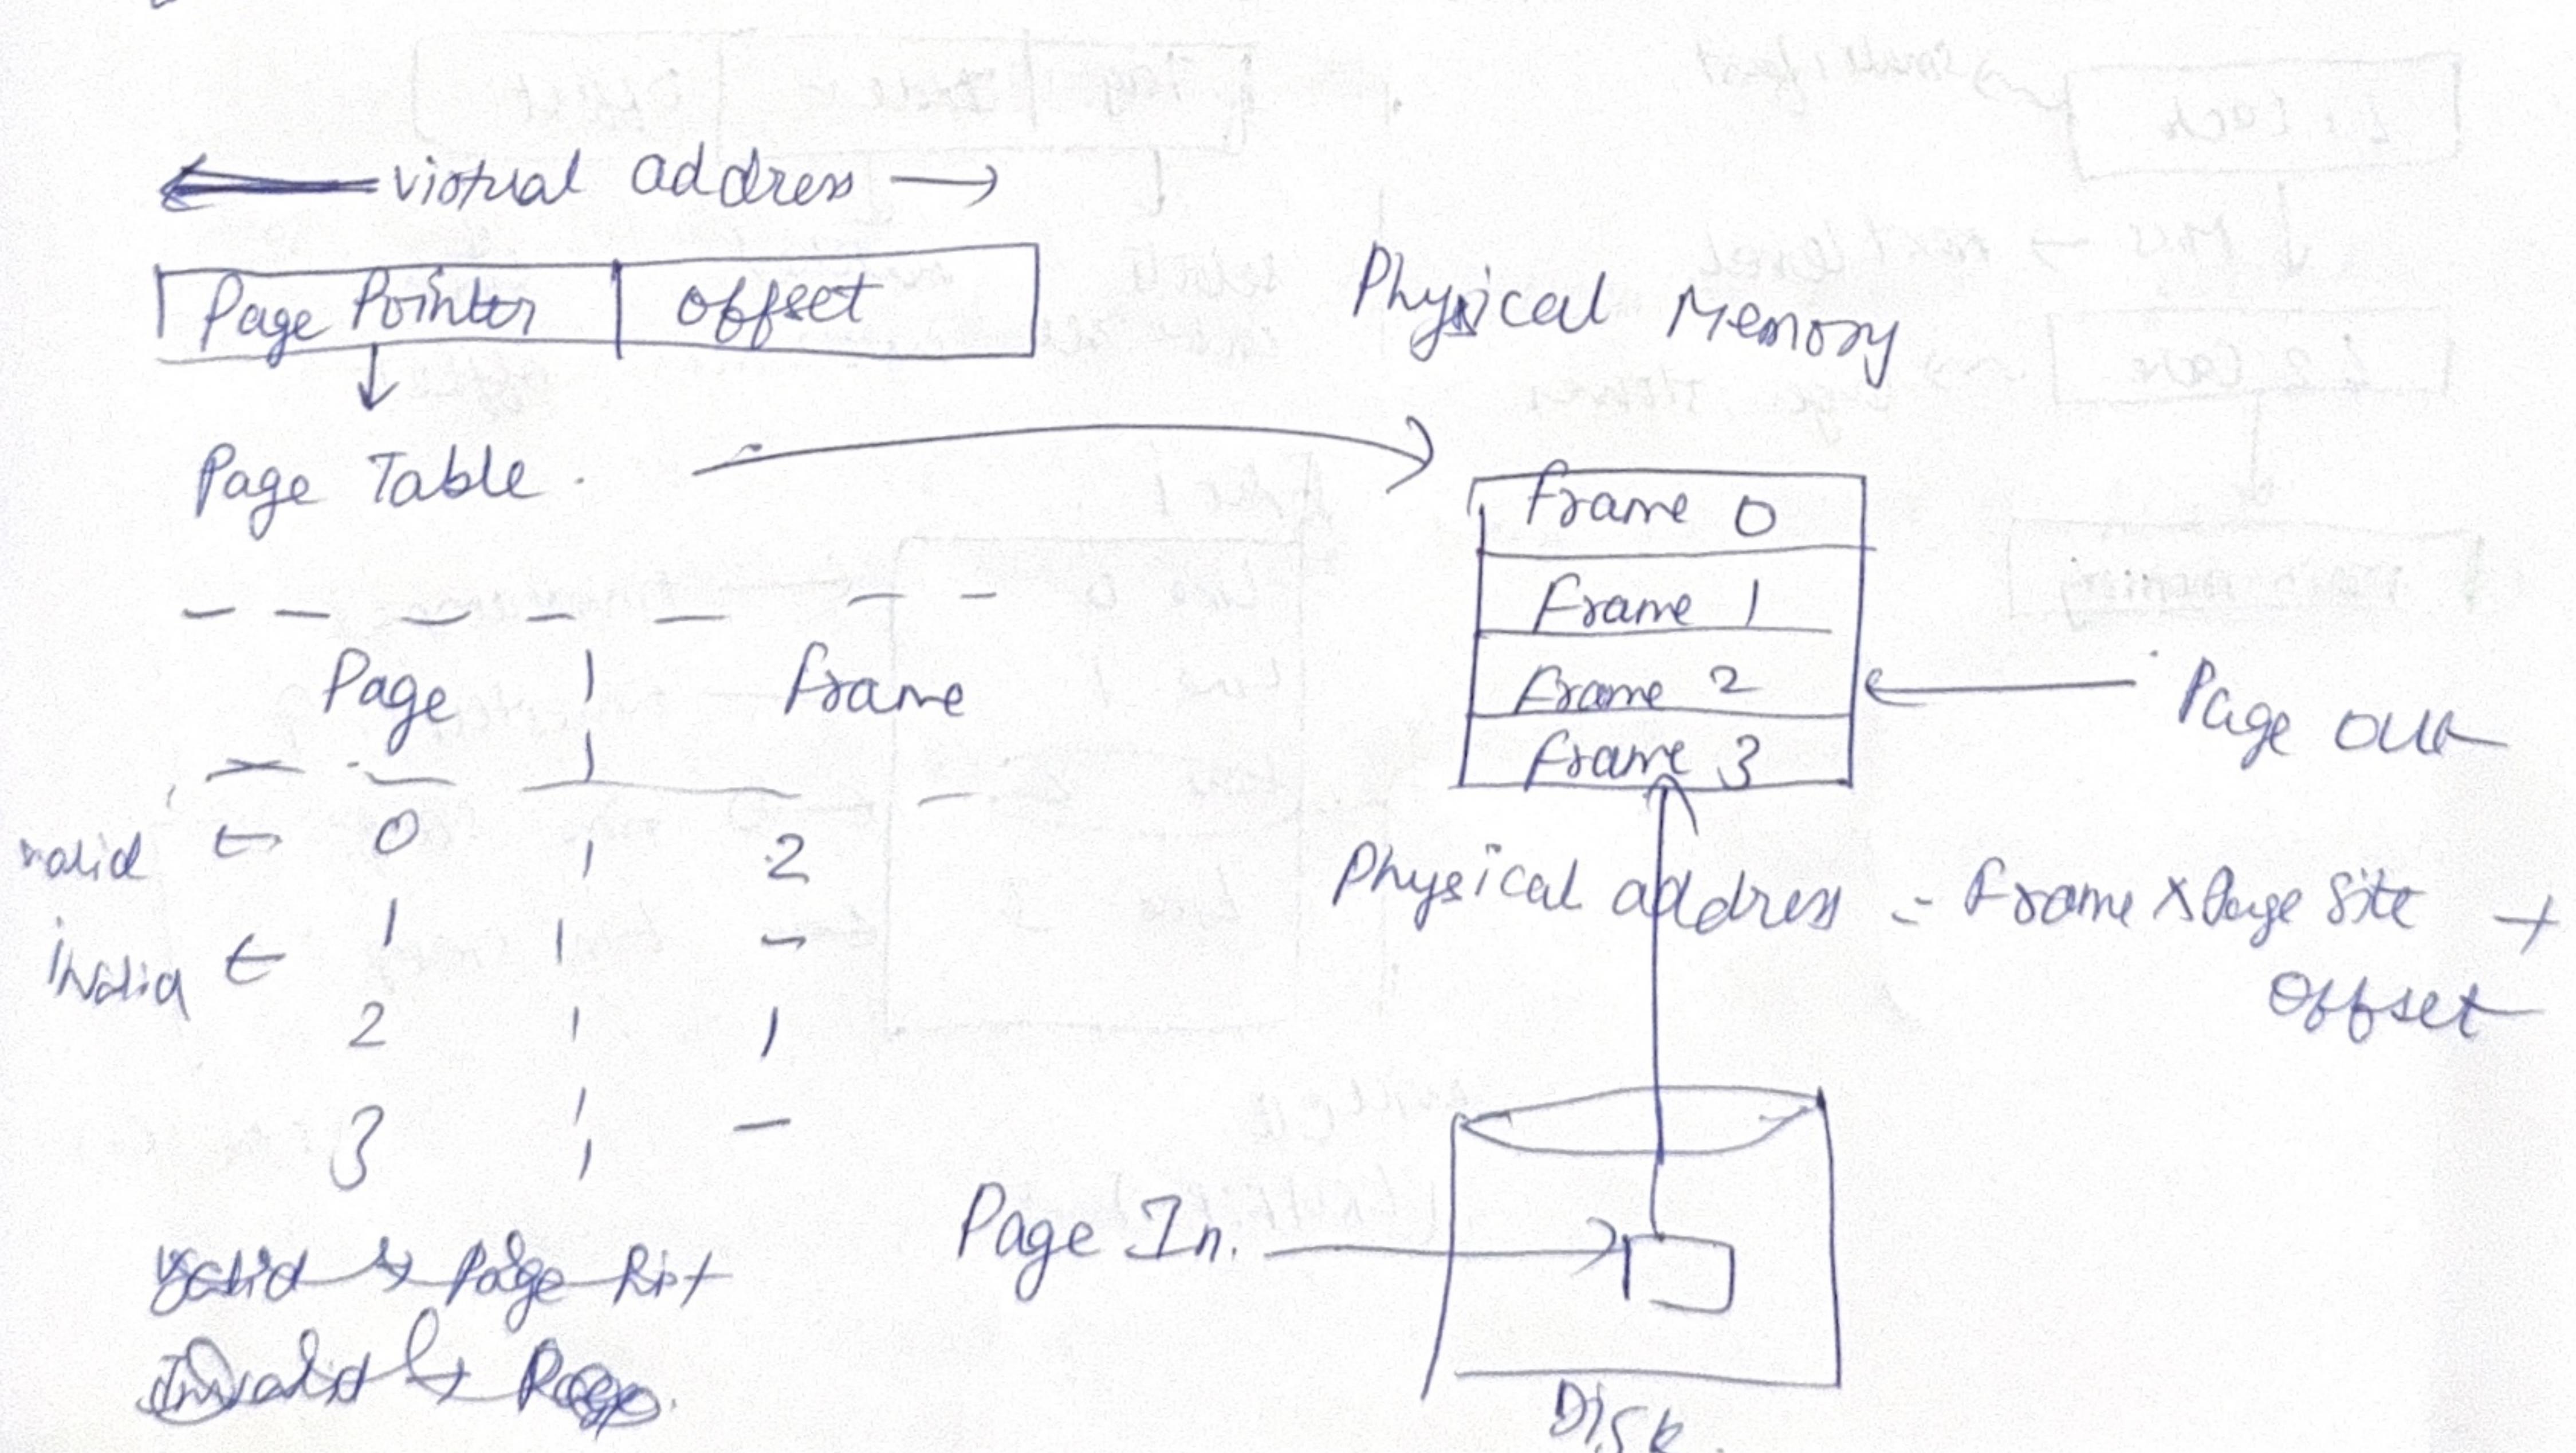
Design Choice: Identity Mapping

To prioritize the fidelity of the Cache and Allocation subsystems, this version of the simulator implements Identity Mapping for virtual memory.

\rightarrow Mechanism: Virtual Address (VA) == Physical Address (PA)

\rightarrow Rationale: This removes the overhead of page tables and Translation Lookaside Buffers (TLB), allowing the simulation to focus on the interplay between variable-size memory segments and cache lines.

\rightarrow Behaviour: All memory requests from the shell are passed directly to the Cache Hierarchy logic without intermediate translation.



- ⑤ Address Translation Flow - since identity mapping is used, the flow is streamlined
- a) User request: malloc(16) returns Virtual Address 0
 - b) Translation : VA 0 maps to physical address 0
 - c) Cache Looking :
 - Address 0 is binary ... 0000 0000.
 - L1 Index 00 (Line 0). Check Line 0.
 - L2 Index 0000 (Line 0). Check Line 0.
 - d) Memory Access : If both cache miss, data is fetched from Physical RAM at index 0.

⑥ Limitations and Implications

- a) No Disk Swap - The system does not simulate hard disk storage or page faults (swapping pages in/out). It assures all memory is resident.
- b) Single Process - The Simulator manages a single global memory space, rather than separate process spaces for multiple processes.
- c) Direct Mapped Cache : While simple, this strategy suffers from conflict misses compared to ~~assoc~~ Set-Association caches.
- d) Buddy System - A buddy allocator was considered but first fit was chosen to allow for arbitrary allocation sizes without strict power-of-two alignment constraints, reducing Internal Fragmentation.